

Singleton

Definition

Class Diagram

Implementation of Singleton Pattern

1. Eager Initialization
2. Lazy Initialization
3. Thread-Safe/Synchronized Initialization
4. Double Locking

Issue with Double-Checked Locking

Issue 1: Instruction Reordering

Issue 2: L1 Caching

The Correct Solution: Using the volatile keyword

▼ Resources

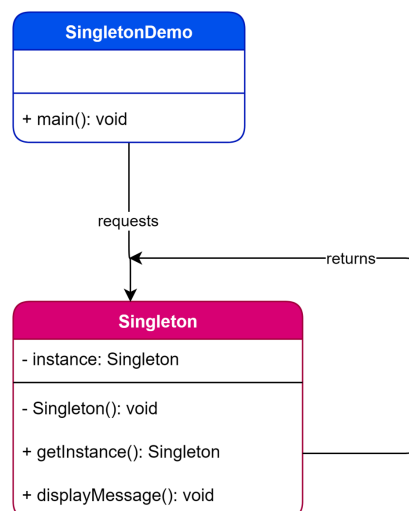
• Video → [27. All Creational Design Patterns | Prototype, Singleton, Factory, AbstractFactory, Builder Pattern](#)

• Video → [28. BUG in Double-Checked Locking of Singleton Pattern & its Fix | Low Level System Design Question](#)

Definition

The Singleton design pattern is used when we have to create only ONE instance of a class. This is useful when exactly one object is needed to perform and coordinate many actions across the system. A Singleton class ensures that there is only one instance of itself, regardless of the number of clients attempting to access it.

Class Diagram



Implementation of Singleton Pattern

There are 4 ways to implement the Singleton Pattern, each with its advantages and disadvantages.

1. Eager Initialization

- One of the simplest implementations of the Singleton Pattern.
- The instance is created when the class is loaded.
- Make the constructor **private** → prevents instantiation by other classes.
- Mark the singleton instance field as:
 - **static** variable → this ensures there's only one class instance shared across all instances of the class.
 - **final** variable → prevents the singleton instance from being reassigned after initialization.

Implementation

```
1 // 1. Eager Initialization - Singleton
2 public class DBConnectionEager {
3
4     // The single instance, created immediately
5     private static final DBConnectionEager instance = new
    DBConnectionEager();
6
7     // The private constructor prevents instantiation
8     private DBConnectionEager() {
9     }
10
11    // Method to return the unique instance of this class
12    public static DBConnectionEager getInstance() {
13        return instance;
14    }
15
16    // Method to display a message
17    public void displayMessage() {
18        System.out.println("Eager Initialization - Singleton - " +
    this);
19    }
20 }
21 // Test Singleton Implementation
22 public class SingletonDemo {
23     public static void main(String[] args) {
24         System.out.println("==== Testing Eager Initialization
    ====");
25         // SingletonObject obj = new SingletonObject(); -->
    Compilation error
26         // Get the unique instance of SingletonObject
27         DBConnectionEager eager1 = DBConnectionEager.getInstance();
28         DBConnectionEager eager2 = DBConnectionEager.getInstance();
29         // Display the message
30         eager1.displayMessage();
31         eager2.displayMessage();
32         // Check if the instances are the same
33         System.out.println("Same instance? " + (eager1 == eager2));
34     }
35 }
```

2. Lazy Initialization

- The instance is created only once when the client requests it.
- Make the constructor **private** → prevents instantiation by other classes.
- It is not thread-safe. When multiple threads access the **getInstance** method for the first time, there is a chance of multiple singleton instances being created.
- Mark the singleton instance field as:
 - **static** variable → this ensures there's only one class instance shared across all instances of the class.

Implementation

```
1 // 2. Lazy Initialization Singleton
2 public class DBConnectionLazy {
3
4     private static DBConnectionLazy instance = null;
5
6     // The private constructor prevents instantiation
7     private DBConnectionLazy() {
8     }
9
10    // Singleton Object is created only when it is required
11    // This method returns the unique instance of this class
12    // Drawback: This implementation is not thread-safe.
13    public static DBConnectionLazy getInstance() {
14        if (instance == null) {
15            instance = new DBConnectionLazy();
16        }
17        return instance;
18    }
19
20    // Method to display a message
21    public void displayMessage() {
22        System.out.println("Lazy Initialization - Singleton - " +
23        this);
24    }
25 }
26 // Test Singleton Implementation
27 public class SingletonDemo {
28     public static void main(String[] args) {
29         System.out.println("===== Testing Lazy Initialization =====");
30         // Get the unique instance of SingletonObject
31         DBConnectionLazy lazy1 = DBConnectionLazy.getInstance();
32         DBConnectionLazy lazy2 = DBConnectionLazy.getInstance();
33         // Display the message
34         lazy1.displayMessage();
35         lazy2.displayMessage();
36         // Check if the instances are the same
37         System.out.println("Same instance? " + (lazy1 == lazy2)); //
38         true
39     }
40 }
```

3. Thread-Safe/Synchronized Initialization

- This is similar to Lazy Initialization
- It is made thread-safe by using the **synchronized** keyword in the **getInstance** method, ensuring only one thread can execute this method at a time.
- This will ensure only one singleton instance is created when multiple threads are invoking the same **getInstance** method.
- But this process of singleton instantiation can be expensive because when a thread enters the synchronized method, it acquires a **lock** on the **class object**, does the job, and releases the lock for other threads to acquire it and execute the method.
 - Imagine 100 threads invoking the synchronized **getInstance** method simultaneously; 99 will have to wait until the first thread finishes its execution.

Implementation

```
1 // 3. Thread Safe Singleton
2 public class DBConnectionThreadSafe {
3
4     // Singleton Object is created only when it is required
```

```

5     private static DBConnectionThreadSafe instance = null;
6
7     // Private Constructor to avoid client applications from using the
    constructor
8     private DBConnectionThreadSafe() {
9     }
10
11    // Thread Safe Method to return the unique instance of this class
12    public static synchronized DBConnectionThreadSafe getInstance() {
13        if (instance == null) {
14            instance = new DBConnectionThreadSafe();
15        }
16        return instance;
17    }
18
19    // Method to display a message
20    public void displayMessage() {
21        System.out.println("Thread Safe Singleton - " + this);
22    }
23 }
24 // Test Singleton Implementation
25 public class SingletonDemo {
26     public static void main(String[] args) {
27         System.out.println("===== Testing Thread Safe =====");
28         // Get the unique instance of SingletonObject
29         DBConnectionThreadSafe threadSafe1 =
    DBConnectionThreadSafe.getInstance();
30         DBConnectionThreadSafe threadSafe2 =
    DBConnectionThreadSafe.getInstance();
31         // Display the message
32         threadSafe1.displayMessage();
33         threadSafe2.displayMessage();
34         // Check if the instances are the same
35         System.out.println("Same instance? " + (threadSafe1 ==
    threadSafe2)); //true
36     }
37 }

```

4. Double Locking

- It is a more optimised version of thread-safe singleton object instantiation.
- This method reduces performance overhead from synchronization(as seen previously) by only synchronizing the block of code when the singleton object is initially created(the first time). Upon instantiation, the other/future threads do not have to enter the synchronized block. Thus, making it faster.
- Widely used in the industry.

Implementation

```

1 // 4. Double Locking Singleton
2 public class DBConnectionDoubleLocking {
3
4     // Double Locking Singleton instance variable
5     private static DBConnectionDoubleLocking instance = null;
6
7     // Private constructor
8     private DBConnectionDoubleLocking() {
9     }
10
11    // Thread Safe Method to return the unique instance of this class
12    public static DBConnectionDoubleLocking getInstance() {
13        if (instance == null) { // first check
14            synchronized (DBConnectionDoubleLocking.class) {
15                if (instance == null) { // second check
16                    instance = new DBConnectionDoubleLocking();
17                }
18            }
19        }
20    }
21 }

```

```

20         return instance;
21     }
22
23     // Method to display a message
24     public void displayMessage() {
25         System.out.println("Double Locking Singleton - " + this);
26     }
27 }
28 // Test Singleton Implementation
29 public class SingletonDemo {
30     public static void main(String[] args) {
31         System.out.println("===== Testing Double Locking =====");
32         // Get the unique instance of SingletonObject
33         DBConnectionDoubleLocking doubleLocking1 =
DBConnectionDoubleLocking.getInstance();
34         DBConnectionDoubleLocking doubleLocking2 =
DBConnectionDoubleLocking.getInstance();
35         // Display the message
36         doubleLocking1.displayMessage();
37         doubleLocking2.displayMessage();
38         // Check if the instances are the same
39         System.out.println("Same instance? " + (doubleLocking1 ==
doubleLocking2)); // true
40
41     }
42 }

```

Issue with Double-Checked Locking

Video explanation → [📺 28. BUG in Double-Checked Locking of Singleton Pattern & its Fix | Low Level System Design Question](#)

Consider the following example with a member variable `portNumber` :

```

1 // Singleton - Double Checked Locking - Issue Demo
2 public class DBConnectionDoubleCheckedLockIssue {
3
4     // Double Locking Singleton instance variable
5     private static DBConnectionDoubleCheckedLockIssue connectionObj =
null;
6     int portNumber;
7
8     // Private constructor
9     private DBConnectionDoubleCheckedLockIssue(int portNumberValue) {
10         portNumber = portNumberValue;
11     }
12
13     // Thread Safe Method to return the unique instance of this class
14     public static DBConnectionDoubleCheckedLockIssue
getConnectionObj() {
15         if (connectionObj == null) { // First check
16             synchronized (DBConnectionDoubleCheckedLockIssue.class) {
17                 if (connectionObj == null) { // Second check
18                     connectionObj = new
DBConnectionDoubleCheckedLockIssue(5567);
19                 }
20             }
21         }
22         return connectionObj;
23     }
24
25     // Method to display a message
26     public void displayMessage() {
27         System.out.println("Singleton - Double Checked Locking - Issue
- " + this);
28     }
29 }

```

```

public static DBConnectionDoubleCheckedLockIssue getConnectionObj() {
    if (connectionObj == null) { // First check
        synchronized (DBConnectionDoubleCheckedLockIssue.class) {
            if (connectionObj == null) { // Second check
                | connectionObj = new DBConnectionDoubleCheckedLockIssue( portNumberValue: 5567);
            }
        }
    }
    return connectionObj;
}

```

The above highlighted step(creation of object) involves the following steps:

1. Allocating memory
2. Initializing the object with all member variables
3. Assigning the reference

Issue 1: Instruction Reordering

The JVM can reorder the object construction steps to achieve better performance. The JVM might **assign the reference before fully initializing the object**, causing other threads to see a **partially constructed object**.

```

public static DBConnectionDoubleCheckedLock getConnectionObj() {
    if (connectionObj == null) { // First check
        synchronized (DBConnectionDoubleLocking2.class) {
            if (connectionObj == null) { // Second check
                L1 memoryPointer = allocateMemory();
                L2 memoryPointer.portNumber= initializeVariable();
                L3 connectionObj = memoryPointer;
            }
        }
        return connectionObj;
    }
}

```

⚠ While T1 is executing the reordered instructions, when T1 is executing L3, let's say T2 performs the first check, it gets the **connectionObject** i.e. **NOT NULL** and the **portNumber** holds the default value(after execution of L2 by T1) and proceeds with performing operations using the partially constructed **connectionObject**, which can result in an error as T1 is yet to initialize the **portNumber** (member variable).


Issue 2: L1 Caching

Without proper synchronization, changes made by one thread may not be visible to other threads because of CPU caching and memory models.

```

public static DBConnectionDoubleCheckedLockIssue getConnectionObj() {
    if (connectionObj == null) { // First check
        synchronized (DBConnectionDoubleCheckedLockIssue.class) {
            if (connectionObj == null) { // Second check
                | connectionObj = new DBConnectionDoubleCheckedLockIssue( portNumberValue: 5567);
            }
        }
    }
    return connectionObj;
}

```



⚠ In a multicore CPU, each CPU core has its own cache; sometimes, cores might not synchronize their caches, leading to inconsistent views of memory. When T1 creates a fully constructed `connectionObject` and saves it in the L1 cache, and if another thread T2 requests a `connectionObject` before the Caches have been synchronised and the changes have been saved in memory, T2 will proceed to create another singleton `connectionObject` instance and perform operations with it. This will lead to multiple singleton instances.

The Correct Solution: Using the `volatile` keyword

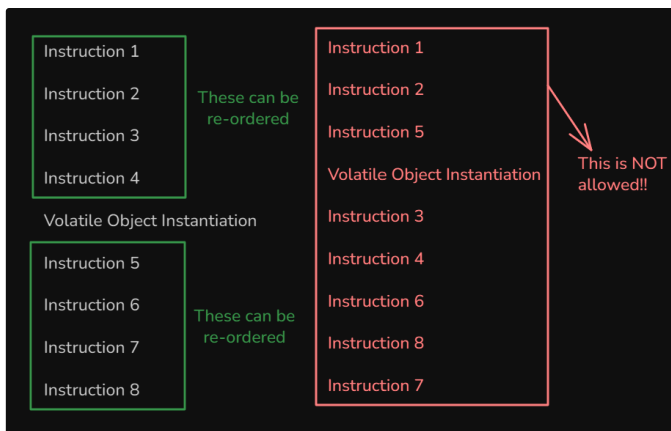
Understanding the `volatile` keyword

When a variable is declared as `volatile`, it instructs the JVM and the compiler to handle that variable in a specific way. This modifier is primarily used in multithreaded programming to ensure memory visibility, consistency of shared variables across different threads, and instruction ordering.

How do `volatile` keyword properties solve the issues with double-checked locking?

- **Memory Visibility Guarantee:** All reads and writes of a `volatile` variable are always performed directly to and from the main memory, and all write operations are immediately visible to all threads. Without `volatile`, threads might cache variables locally and not see updates from other threads. [\[Solves Issue 2\]](#)
- **Instruction Re-ordering:** The `volatile` keyword provides guarantees to prevent specific types of instruction reordering, as illustrated below. This prevention is accomplished by establishing "happens-before" guarantees, which effectively create memory barriers or fences. These barriers stop the compiler and CPU from reordering instructions around `volatile` reads or writes, even if such reordering would enhance performance in a single-threaded context.

[\[Solves Issue 1\]](#)



```
public static DBConnectionDoubleCheckedLock getConnectionObj() {
    if (connectionObj == null) { // First check
        synchronized (DBConnectionDoubleLocking2.class) {
            if (connectionObj == null) { // Second check
                memoryPointer = allocateMemory();
                memoryPointer.portNumber = initializeVariable();
                connectionObj = memoryPointer;
            }
        }
    }
    return connectionObj;
}
```

only these two instructions are allowed to be reordered (but, in this cases its not logical to re-order these two)

volatile object instruction

Implementation: Using the volatile keyword

```
1 // Singleton - Double Checked Locking - Fix Demo
2 public class DBConnectionDoubleCheckedLockFix {
3
4     // volatile keyword is used to ensure that the value of the
5     // variable
6     // is fetched from the memory every time.
7     private static volatile DBConnectionDoubleCheckedLockFix
8     connectionObj = null;
9
10    int portNumber;
11
12    private DBConnectionDoubleCheckedLockFix(int portNumberValue) {
13        portNumber = portNumberValue;
14    }
15
16    // Thread Safe Method to return the unique instance of this class
17    public static DBConnectionDoubleCheckedLockFix
18    getConnectionObj(int portNumberValue) {
19        if(connectionObj == null) {
20            synchronized(DBConnectionDoubleCheckedLockFix.class) {
21                if(connectionObj == null) {
22                    connectionObj = new
23                    DBConnectionDoubleCheckedLockFix(portNumberValue);
24                }
25            }
26        }
27        return connectionObj;
28    }
29
30    // Method to display a message
31    public void displayMessage() {
32        System.out.println("Singleton - Double Checked Locking - Fix -
33        " + this);
34    }
35 }
```